

CSC 212: Data Structures and Abstractions

06: Dynamic Arrays (part 2)

Prof. Marco Alvarez

Department of Computer Science and Statistics
University of Rhode Island

Fall 2026



Where we stand?

- Arrays give $\Theta(1)$ indexing, because memory is contiguous
- Inserting or removing in the middle costs $\Theta(n)$, because data must move
- Grow by constant costs $\Theta(n^2)$ for n appends
- Open question:
 - ✓ is there a growth rule where n appends cost only $\Theta(n)$?

2

Grow by doubling

- When full, grow to: $2 \times \text{capacity}$ (starting from 1)
 - ✓ copies get bigger, but they get rarer


capacity: 1 $\rightarrow$ 2 $\rightarrow$ 4 $\rightarrow$ 8 $\rightarrow$ 16 $\rightarrow$ 32 $\rightarrow$...

push 1:	[1]	grow to 1
push 2:	[1 2]	full: copy 1, grow to 2
push 3:	[1 2 3 _]	full: copy 2, grow to 4
push 4:	[1 2 3 4]	room: no copy
push 5:	[1 2 3 4 5 _ _]	full: copy 4, grow to 8
push 6-8:	[1 2 3 4 5 6 7 8]	room: no copies
push 9:	[1 2 ... 9 _ _ _ _ _]	full: copy 8, grow to 16

3

Cost of each push_back

- Element writes for each of the first 16 appends
 - ✓ copies + the write itself
 - ✓ total for 16 appends: 31 writes, fewer than **2 per append**



4

Counting the copies

- Copies happen when the size reaches $1, 2, 4, \dots, 2^k$, the powers of 2 below n
 - a geometric series: $\text{copies} = 1 + 2 + 4 + \dots + 2^k = 2^{k+1} - 1$
 - since $2^k < n$, we get $2^{k+1} < 2n$, so $\text{copies} < 2n$

- Count element writes for n appends, starting **empty**
 - ignore cost of allocating/deallocating memory

append	copies	writes
1	0	1
2	1	1
3	2	1
4	—	1
5	4	1
6	—	1
7	—	1
8	—	1
9	8	1
10	—	1
...	...	...
n-1	—	1
n	—	1

$$\begin{aligned}
 T(n) &= n + \sum_{i=0}^{\log_2 n - 1} 2^i \\
 &= n + \left(\frac{2^{\log_2 n} - 1}{2 - 1} \right) \\
 &= n + (n - 1) \\
 &= \Theta(n) \quad \text{cost of adding } n \text{ elements}
 \end{aligned}$$

assume n is a power of 2

The **amortized cost** of inserting an element is $\Theta(1)$ and any sequence of n insertions takes at most $\Theta(n)$ time in total.

5

Amortized cost

- Amortized cost is the **average cost per operation**
 - over a sequence of operations, even if some individual operations are expensive
 - simple way to analyze cost where occasional costly operations are offset by cheaper ones
- Dynamic arrays with doubling
 - cost is $\Theta(1)$ when there is space, $\Theta(n)$ when growing
 - over n appends, the total cost is $\Theta(n)$
 - the amortized cost per insertion is $\Theta(1)$

6

Q6-1

- A dynamic array doubles when full. Which statement is **false**?

- A single `push_back` is $O(n)$ in the worst case
- n calls to `push_back` take $\Theta(n)$ in total
- n calls to `push_back` take $O(n^2)$ in total
- n calls to `push_back` take $\Theta(n^2)$ in total
- None of the above


7

Answer

- D) n calls to `push_back` take $\Theta(n^2)$ in total

- the others are true
 - A) a call that triggers a grow copies all n elements
 - B) the tight bound
 - C) true, but **loose**, O is only an upper bound

8



Grow by factor

- Any factor greater than 1 works
 - when full, grow to: $r \times \text{capacity}$, for any $r > 1$

$$\text{copies} < n \frac{r}{r-1} = \Theta(n)$$

factor	copies, at most	unused slots, after a
2	2n	up to 50%
1.5	3n	up to 33%
1.123	9n	up to 11%

smaller factor: less waste, more copies

10

Growth factors in real libraries

- C++ (`std::vector`)
 - implementation-defined (not in the C++ standard)
 - grow by 1.5 (MS Visual C++) or 2.0 (g++/clang)
- Java (`ArrayList`)
 - grow by 1.5 in OpenJDK
- Python (`List`)
 - grow by ~ 1.125
- Rust (`std::vec::Vec`)
 - typically grow by 2

https://en.wikipedia.org/wiki/Dynamic_array

11

Q6-2

- Which growth rule makes n appends cost $\Theta(n^2)$?

- Grow to 1.01 x capacity
- Grow to 1.5 x capacity
- Grow to capacity + 10,000
- Grow to 2 x capacity
- None of the above

12

Answer

- C) Grow to capacity + 10,000
- adding a constant is $\Theta(n^2)$ cost, whatever the constant

13

Shrinking the array

- After many `pop_back` calls, most of the capacity may be empty
 - should shrink
- First idea, halve the capacity when array is half full
 - **worst-case** when array is full and alternate between adding and removing
 - each alternating operation would require resizing the array
- More efficient resizing
 - half the capacity when the array is **one-quarter** full
- In practice ...
 - most standard implementations do not automatically shrink capacity
 - avoids performance penalties from frequent resizing
 - instead, they provide explicit operations like `shrink_to_fit()` that allow the programmer to request size reduction when deemed necessary

14

Append and indexing

```
void DynArray::push_back(int val) {
    if (m_size == m_capacity) {
        grow(m_capacity == 0 ? 1 : 2 * m_capacity); // double
        m_data[m_size++] = val;
    }

    int& DynArray::operator[](size_t i) {
        return m_data[i]; // no check: fast
    }

    int& DynArray::at(size_t i) {
        if (i >= m_size) // checked: safe
            throw std::out_of_range("DynArray::at");
        return m_data[i];
    }
}
```


- `operator[]` trusts the caller, at checks (same behavior as `std::vector`)
- Returning `int&` lets the caller write: `a[3] = 7`

15

The copy problem

- C++ writes a copy constructor for us. It copies each **member**, including the pointer

```
{
    DynArray a;
    a.push_back(1);
    DynArray b = a; // default copy: copies m_data, the address
} // end of scope: ~b frees the block, then ~a frees it again
```



One block, two owners. A **shallow copy** copies the pointer, a **deep copy** copies the data.

16

Rule of three, rule of five

- A class that owns memory **needs all three, or none**
 - ✓ **destructor**: frees the block
 - ✓ **copy constructor**: `DynArray b = a;` makes a deep copy
 - ✓ **copy assignment**: `b = a;` frees b's block, then makes a deep copy
- Rule of five adds move operations: take over a temporary's block instead of copying it

```
DynArray::DynArray(const DynArray& other) // perform deep copy
: m_data(other.m_capacity ? new int[other.m_capacity] : nullptr),
  m_size(other.m_size), m_capacity(other.m_capacity) {
    for (size_t i = 0; i < m_size; ++i)
        m_data[i] = other.m_data[i];
}

DynArray& DynArray::operator=(const DynArray& other) {
    if (this != &other) {
        DynArray copy(other); // deep copy first
        std::swap(m_data, copy.m_data); // take the copy's block
        std::swap(m_size, copy.m_size);
        std::swap(m_capacity, copy.m_capacity);
    } // copy's destructor frees our old block
    return *this;
}
```

17

Q6-3

- An AI assistant wrote this class. It compiles without warnings. What happens at `b = a`?

- A. The copy constructor runs, so `b` gets a deep copy
- B. The default copy assignment runs: `b` shares `a`'s block, `b`'s old block leaks, and the block is freed twice
- C. It does not compile, because `operator=` is missing
- D. `b = a` moves the data from `a` into `b`
- E. None of the above

```
class DynArray {
private:
    int* m_data;
    size_t m_size;
    size_t m_capacity;

public:
    DynArray() : m_data(nullptr), m_size(0), m_capacity(0) {}
    DynArray(int val) : m_data(new int[val]), m_size(val), m_capacity(val) {}
    DynArray(const DynArray& other) : m_data(new int[other.m_capacity]), m_size(other.m_size), m_capacity(other.m_capacity) {
        for (size_t i = 0; i < m_size; ++i)
            m_data[i] = other.m_data[i];
    }
    void push_back(int val) { // doubles when full, as before
    };

    int main() {
        DynArray a;
        a.push_back(1);
        DynArray b;
        b = a;
    }
}
```

18

Answer

- B) The default copy assignment runs: `b` shares `a`'s block, `b`'s old block leaks, and the block is freed twice
- ✓ `DynArray b = a;` is **construction** (copy constructor is called)
- ✓ `b = a;` on an existing object is **assignment**: not written in the code, so C++ copies the members
- ✓ result: a leak (`b`'s old block) and a double free (at the end of `main`)

19

std::vector from C++ STL

```
#include <iostream>
#include <vector>

int main() {
    std::vector<int> v;
    size_t last = v.capacity();
    for (int i = 0; i < 1000; ++i) {
        v.push_back(i);
        if (v.capacity() != last) {
            last = v.capacity();
            std::cout << "size " << v.size() << " -> capacity " << last << '\n';
        }
    }
}
```

What is the output?

```
size 1 -> capacity 1
size 2 -> capacity 2
size 3 -> capacity 4
size 5 -> capacity 8
size 9 -> capacity 16
...
size 513 -> capacity 1024
```

20

Iterator invalidation

```
#include <iostream>
#include <vector>

int main() {
    std::vector<int> v = {1, 2, 3};
    int& first = v[0];
    v.push_back(4); // may reallocate: first now dangles
    std::cout << first << '\n'; // undefined behavior
}
```

What is the output?

After any call that may grow a vector (push_back, insert, resize, reserve), pointers, references, and iterators to its elements may be invalid. Keep indices instead.

<https://en.cppreference.com/w/cpp/container/vector>

21

Q6-4

- Fill in the table for a dynamic array that doubles when full and never shrinks. Use Θ

Operation	Best case	Worst case
Append 1 element		
Remove 1 element from the end		
Insert 1 element at index idx		
Remove 1 element at index idx		
Read element at index idx		
Write (update) element at index idx		

22

Q6-5

- Run this code, which loop is $\Theta(n^2)$?

```
import time
n = 100000

start = time.time()
array = []
for i in range(n):
    array.append('s')
print(time.time() - start)

start = time.time()
array = []
for i in range(n):
    array = array + ['s']
print(time.time() - start)
```

- A. The first: `append` copies the list every time
B. The second: `array + ['s']` builds a brand-new list every time
C. Both: Python lists are linked lists
D. Neither: both are $\Theta(n)$ in total
E. None of the above

23

Answer

- B) The second: `array + ['s']` builds a brand-new list every time

- ✓ `array + ['s']` allocates a new list and copies all k elements into it: that is **grow by one**
- ✓ `append` grows the list in place, with spare capacity

- ✓ run (Python 3.9, the same machine as before)

- for $n = 100,000$, `append` takes 0.0075 s and the grow by one method takes 11.7 s (about 1,500x slower)

- ✓ Python lists are dynamic arrays

24

Designing a dynamic array class in C++

```
#include <cstdint>
//
// This simplified implementation uses int. A production-quality dynamic array
// would use C++ templates to support arbitrary types, as std::vector<T> does.
// It also would implement the Rule of Three/Five for proper memory management.
//
class DynamicArray {
private:
    int *_m_arr;           // pointer to the (internal) array
    size_t _m_capacity;    // total number of elements that can be stored
    size_t _m_size;        // number of elements currently stored

public:
    DynamicArray();        // constructor (allocates memory)
    ~DynamicArray();       // destructor (frees memory)
    void push_back(int val); // add an element to the end
    void pop_back();        // remove the last element
    void insert(int val, size_t idx); // insert an element at a specific index
    void erase(size_t idx); // remove an element at a specific index
    void resize(size_t len); // change the capacity of the array
    size_t size() const;    // return the number of elements
    size_t capacity() const; // return the capacity
    bool empty() const;    // check if the array is empty
    void clear();          // remove all elements, maintaining the capacity

    // additional methods can be added here
};
```

A class definition specifies the **data members** and **member functions** of the class. The data members are the attributes of the class, and the member functions are the operations that can be performed on the data members. The class definition is a blueprint for creating objects of the class.